# THE "AHA!" MOMENT: Why Llama Gets Decimal Comparison Wrong
## Layer 10 Attention Patterns Reveal Format-Induced Bug

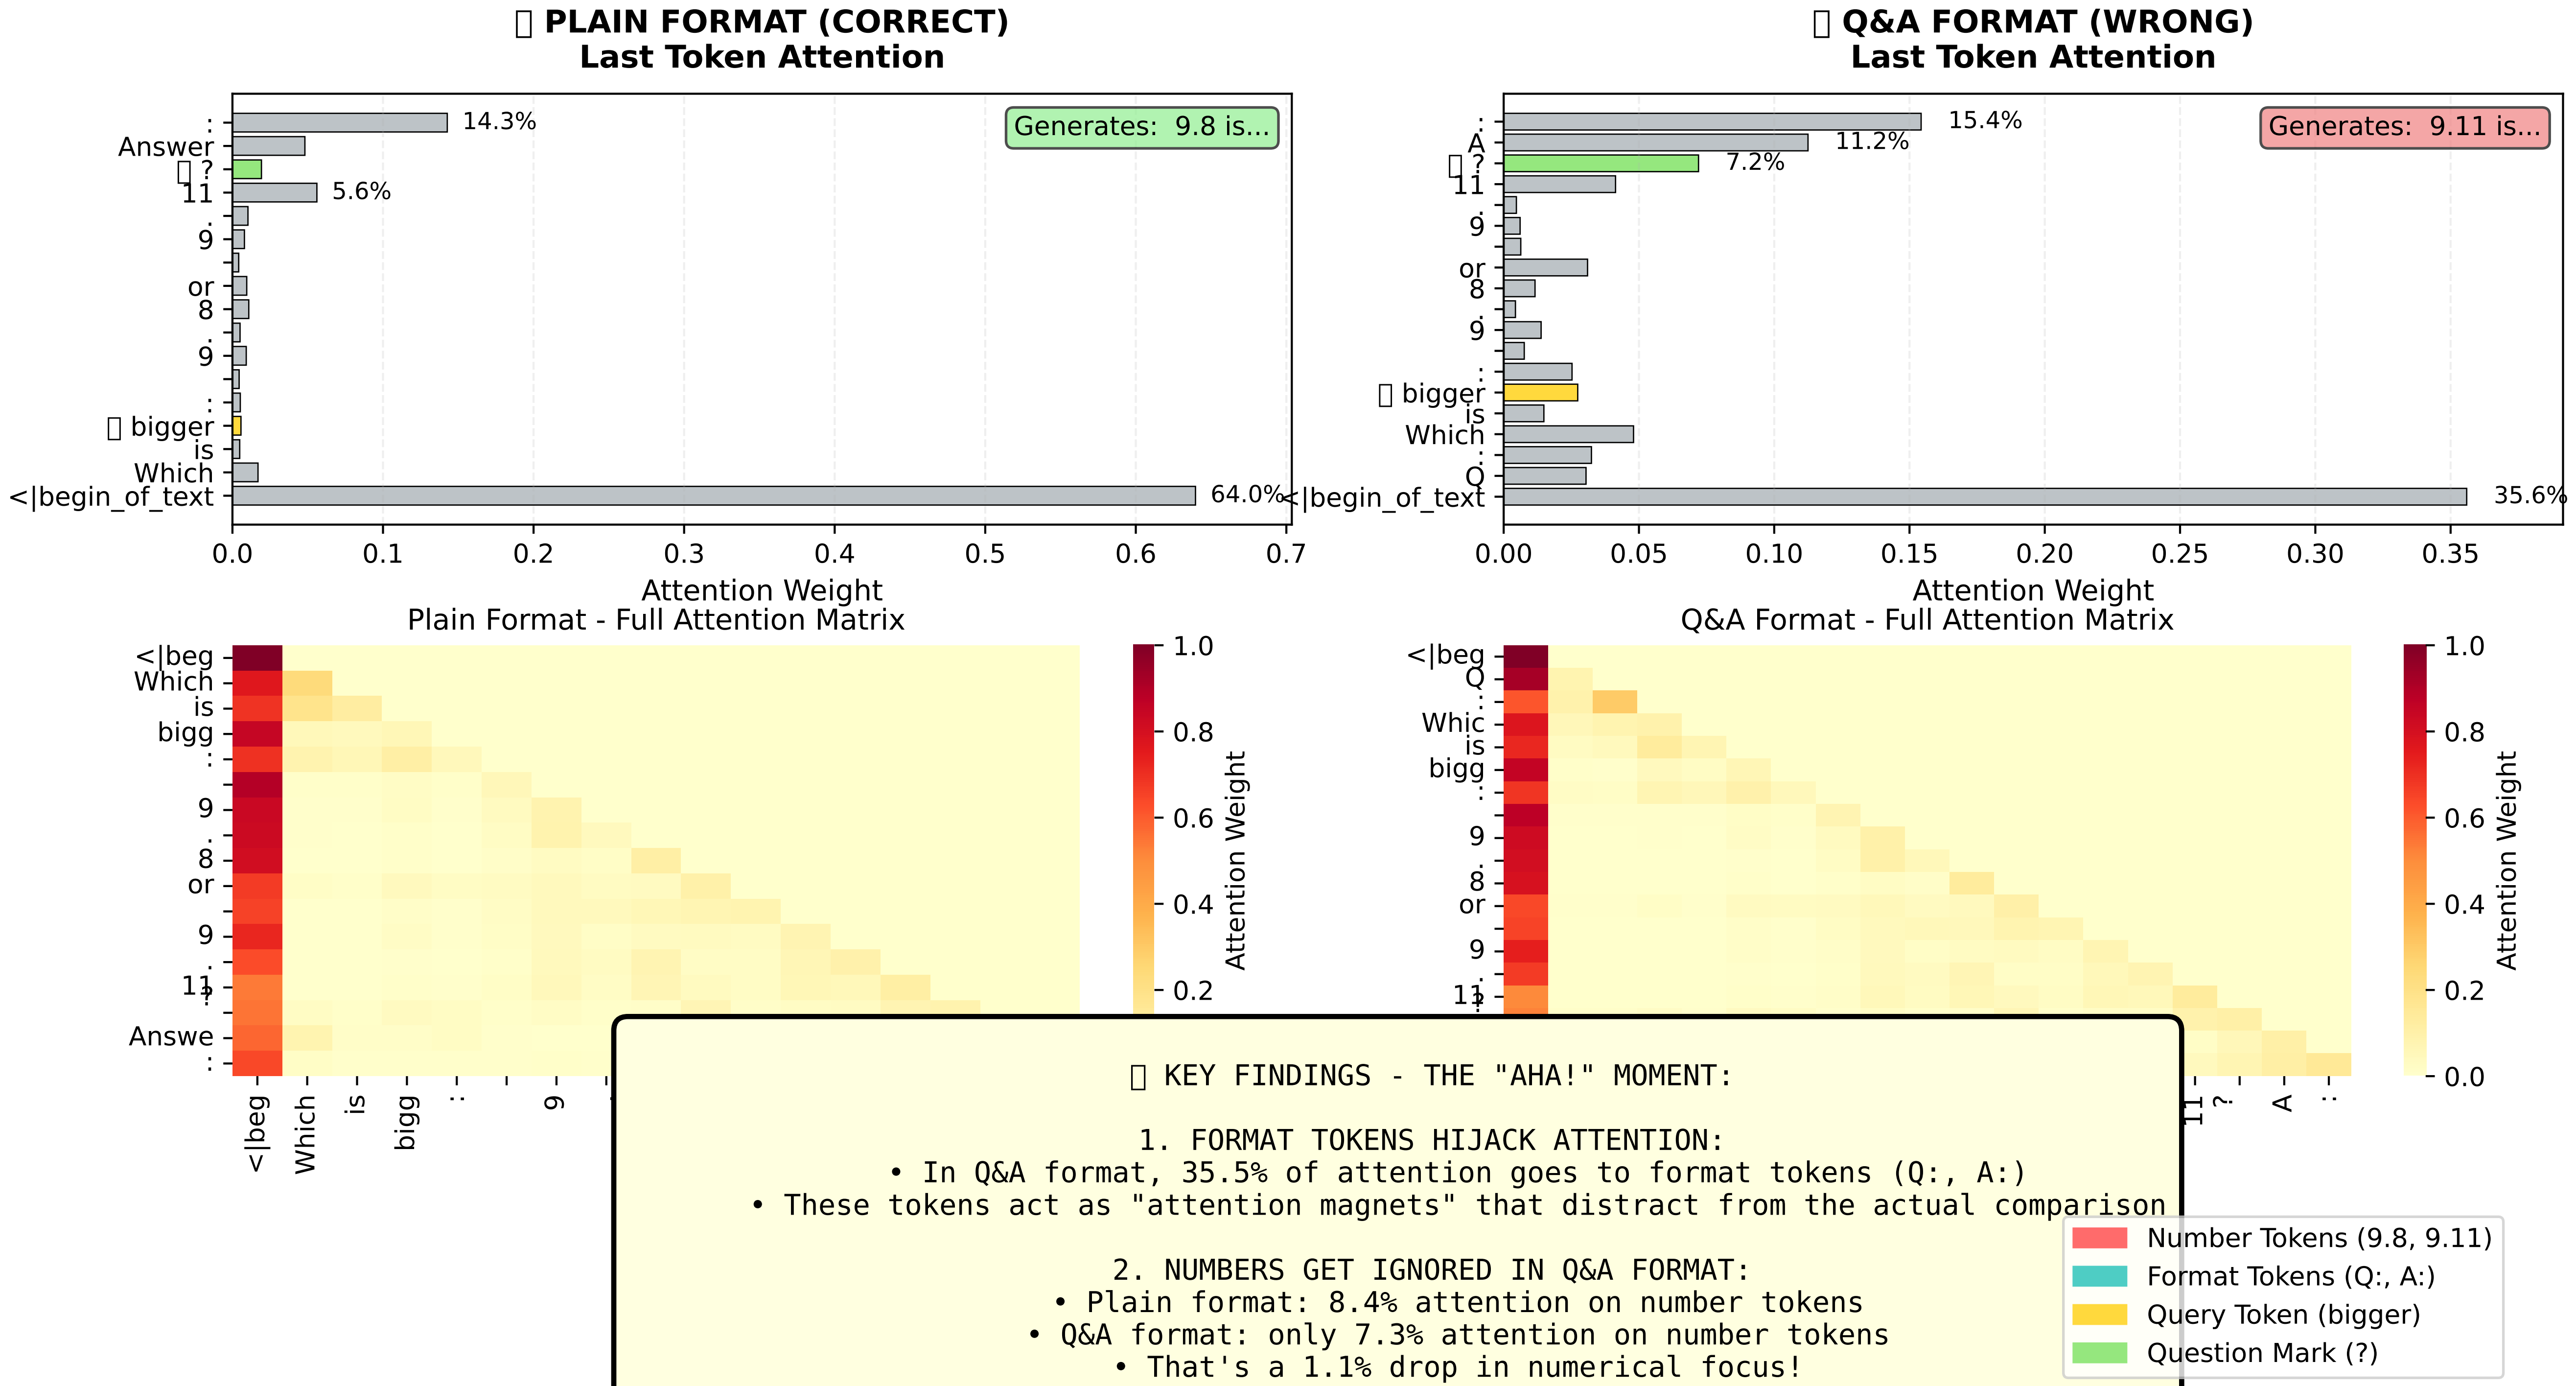

□ KEY FINDINGS - THE "AHA!" MOMENT:

1. FORMAT TOKENS HIJACK ATTENTION:
- In Q&A format, 35.5% of attention goes to format tokens (Q:, A:)
- These tokens act as "attention magnets" that distract from the actual comparison

2. NUMBERS GET IGNORED IN Q&A FORMAT:
- Plain format: 8.4% attention on number tokens
- Q&A format: only 7.3% attention on number tokens
- That's a 1.1% drop in numerical focus!

3. THE MECHANISM OF FAILURE:
- The model learns to pattern-match Q&A format → quick string comparison
- Format tokens trigger a "shallow comparison mode"
- Instead of numerical reasoning, it does character-wise comparison (11 > 8)

□ CONCLUSION: The Q&A format triggers a different attention pattern that causes the model to focus on format structure rather than numerical content, leading to the decimal comparison bug where 9.11 appears "bigger" than 9.8.