



Ablating 8 Hijacker Neurons from Chat Template Circuit:
L7/N1978, L13/N10352, L14/N2451 + 5 more

Ablation Parameter Sweep: Bug Rate vs Coherence Llama-3.1-8B-Instruct



Example outputs: Ablation 0: "9.11 is bigger than 9.8" → Ablation -4.5: "The number 100 is not mentioned..."